

# Kyle Ton

714-757-3573 | [tonkyle06@gmail.com](mailto:tonkyle06@gmail.com) | [www.kylehton.com](http://www.kylehton.com) | [in/kylehton](https://in/kylehton)

## EDUCATION

### University of California, San Diego

June 2027

*Bachelor's in Mathematics and Computer Science*

*La Jolla, CA*

- Relevant Coursework: Object-Oriented Programming, Computer Architecture and Assembly, Data Structures and Algorithms, Discrete Mathematics, Linear Algebra, Differential Equations, Single and Multi-Variable Calculus

## EXPERIENCE

### Full-Stack Software Engineering Intern

June 2025 – Aug. 2025

*ForOurLastNames*

*Remote, CA*

- Containerized and deployed a **Node.js SPA** via **AWS Lambda (Docker)**, serving static content through **S3 and CloudFront** and backend APIs through Lambda to enable near-serverless scalability at sub-\$30 monthly cost.
- Integrated **Supabase Auth and Stripe Billing** with verified webhooks, secure **Row-Level Security (RLS)** policies, and automated 2-week trial flows, establishing a subscription-ready authentication pipeline.
- Configured **Prometheus and Grafana monitoring** using CloudWatch exporter to visualize Lambda invocation metrics, informing capacity planning and resource tuning for an expected 5,000-user scale.

### Full-Stack AI Engineering Intern

Oct. 2024 – Dec. 2024

*PM Accelerator*

*Remote, CA*

- Developed a **FastAPI + React** application powering an AI resume optimizer, collaborating with AI engineers to enhance ATS keyword alignment across 100 test resumes with a **15% average increase**.
- Deployed asynchronous **Celery workers** with a **Redis broker** on Heroku to parallelize resume processing, cutting latency by **50% (20s to 10s)** during prototype testing.
- Created and indexed **MongoDB Atlas** documents linked to Google OAuth IDs, enforcing user-level isolation and safe concurrent editing for 100+ beta users.

### Software Engineering Intern

Aug. 2024 – Oct. 2024

*Trace*

*Garden Grove, CA*

- Automated firmware version tracking and remote diagnostics for **1,000+ cameras** using **Python + Paramiko SSH**, improving documentation accuracy and deployment traceability.
- Developed remote control routines (fan, power toggles, component tests) to isolate hardware issues, reducing troubleshooting time from **2.5 to 1.5 hours** during testing cycles.
- Integrated automated logging into shared **Google Sheets dashboards**, streamlining QA communication and bi-weekly hardware audits.

## PROJECTS

### MatchPointAI | *NumPy, Pandas, Scikit-learn, FastAPI, TypeScript, Next.js, PostgreSQL*

June 2025

- Engineered 12+ player and match level features from a 389,000+ row ATP historical dataset (1968 to 2024) by **pre-processing performance metrics in parallel**, enabling efficient **feature engineering** with 25-match rolling statistics
- Implemented **5-KNN imputation** on 10,000 sample rows to fill missing data, stratified by era and player rank, and **optimized 7+ hyperparameters**, achieving **68%** randomized search cross-validated accuracy on an XGBoost model

### Git-Lint | *Python, FastAPI, Pinecone, AWS Lambda & S3, Docker, OpenAI Agents SDK*

May 2025

- Engineered an AI-powered code reviewer using **FastAPI** and **OpenAI's Agents SDK**, creating an **orchestrative agentic system** capable of context-aware file-by-file analysis and iterative feedback, automating pull request code review suggestions
- Integrated a **RAG-based context retrieval system** using **Pinecone vector search** and **AWS S3** for over **200 code chunks**, enabling access to relevant file embeddings for repo-specific memory across multiple agent steps during code reviews

### StockSense | *React, Java, Spring Boot, PostgreSQL, AWS S3, Docker, Redis*

March 2025

- Built a real-time stock market platform combining **React-based UI** with a Spring Boot backend, delivering live prices, sentiment-tagged news, and storing it in **AWS S3**, with implemented CI/CD via **GitHub Actions**
- Enabled data-driven investment insights by **web-scraping recent and relevant stock news** with yfinance API, analyzing media sentiment with **HuggingFace's FinBERT** model, providing numerical evaluation per-article

## TECHNICAL SKILLS

**Programming Languages:** Java, Python, SQL, JavaScript/TypeScript, HTML/CSS

**Frameworks/Databases:** React.js, Next.js, Spring Boot, FastAPI, MongoDB, PostgreSQL

**Tools/Platforms:** Git, Postman, Docker, Redis, Github Actions, AWS

**Machine Learning:** NumPy, Pandas, Matplotlib, Seaborn, Scikit-learn